Assignment 4

Deadline: 04.06.2021

The questions that includes copy paste from internet (or anywhere) will get 0 points, even if answer is right.

1- What are the advantages and disadvantages of B+ tree for using it for indexing in databases?
2- What are the advantages and disadvantages of hashing for using it for indexing in databases?
3- Indices speed query processing, but it is usually a bad idea to create indices on every attribute, and every combination of attributes, that are potential search keys. Explain why.
4- Is it possible in general to have two clustering indices on the same relation for different search keys? Explain your answer.
5- Suppose you have a relation r with n_r tuples on which a secondary B^+-tree is to be constructed.
    a. Give a formula for the cost of building the B^+-tree index by inserting one record at a time. Assume each block will hold an average of f entries and that all levels of the tree above the leaf are in memory.
    b. Assuming a random disk access takes 10 milliseconds, what is the cost of index construction on a relation with 10 million records?
6- Suppose there is a relation $r(A,B,C)$, with a B^+-tree index with search key (A, B).
    a. What is the worst-case cost of finding records satisfying $10 < A < 50$ using this index, in terms of the number of records retrieved n_1 and the height h of the tree?
    b. What is the worst-case cost of finding records satisfying $10 < A < 50 \wedge 5 < B < 10$ using this index, in terms of the number of records n_2 that satisfy this selection, as well as n_1 and h defined above?
7- Given the relations listed below,

branch (*branch_name*, *branch_city, assets*)
customer (*customer_name*, *customer_street, cust omer_city*)
loan (*loan_number*, *branch_name, amount*)
borrower (*customer_name*, *loan_number*)
account (*account_number*, *branch_name, balance*)
depositor (*customer_name*, *account_number*)

Let us define a view *branch cust* as follows:

create view *branch cust* **as**
select *branch name, customer name*
from *depositor, account*
where *depositor.account number = account.account number*

a- Suppose that the view is *materialized*; that is, the view is computed and stored. Write triggers to *maintain* the view, that is, to keep it up-to-date on insertions to *depositor* or *account*. It is not necessary to handle deletions or updates. Note that, for simplicity, we have not required the elimination of duplicates.

b- Write an SQL trigger to carry out the following action: On **delete** of an account, for each customer-owner of the account, check if the owner has any remaining accounts, and if she does not, delete her from the *depositor* relation.

8- Give characteristics of NoSQL. What is the difference between SQL and NoSQL?